

## PIPES

Create pipe (unidirec comm buff w/ 2 file descriptors fd[0] read & fd[1] write)

```
#include <unistd.h>
int pipe(int fd[2]);
```

Data write & read FIFO base. No external/permanent name; only accessed via 2 fds.

Pipe can only be used by process that created it & its descendants.

close(fd) closes a file descriptor

dup(newfd) duplicates fd, dup2(newfd, oldfd) copies (more like alias)

Read:	Write:
Not nec atomic; may read less bytes	Atomic for at most PIPE_BUF bytes (512, 4k, 64k)
Blocking: if no data, write fd still opens.	Blocking: if buffer full & read fd open.
If empty & all fd for write closes; read sees eof, returns 0	When all fd to read closed, causes SIGPIPE sig for calling

Pros: simple, flexible, efficient comms

Cons: no way to open already existing pipe: impossible for 2 arbitrary processes to share same pipe (unless created by common ancestor)

## IPC: UNIX Shared Mem

Parent & child processes run in separate addr spaces.

Shared mem segment: piece of mem that can be allocated & attached to addr space; process w/ this mem seg attached will have access to it but race conditions can occur

**Procedure:** find a key (unix uses key to identify shared mem segments) -> shmget() allocates a shared mem -> shmat() attach shared mem to addr space -> shmdt() detach mem from addr space -> shmctl() deallocate shared mem

**Keys:** global; If other processes know key, can access shared mem

Can ask sys to provide private key using IPC\_PRIVATE

DIY:

```
key_t SomeKey;
SomeKey = 1234;
Autogenerate:
key_t = ftok(char *path, int ID);
```

**Asking for Shared Memory:**

Use shmget() to request a shared mem (returns shared mem ID):

```
shm_id = shmget(key_t key, int size, int flag)
```

Flag for our purpose either 0666 (rw) or IPC\_CREAT|0666

Include following:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

## Attaching Shared Memory:

Use shmat() to attach existing shared mem to addr space (ret void pointer to mem):

```
shm_ptr = shmat(int shm_id, char *ptr, int flag);
shm_id ID from shmget(), use NULL for ptr, flag = 0
```

**Detaching & Removing Shared Memory:**

To detach shared memory, use shmdt(shm\_ptr);

shm\_ptr is pointer returned by shmat().

After shared mem *detached*, it's still there (can be reattached to use again. *Removing* a shared mem will make it stop existing)

To remove shared memory, use shmctl(shm\_ID, IPC\_RMID, NULL);

shm\_ID is shared mem ID returned by shmget().

## INTERPROCESS COMMUNICATION

Types: Message passing (Blocking/non-blocking, Datagrams, virtual circuits, streams,

Remote Procedure Calls (RPC)), Shared Memory

**Shared Memory:** >=2 processes share part of their addr space

Adv: fast & easy to use (but occur access to shared data can cause issues, and must sync access to shared data)

Disadv: senders & receivers must be on same machine, less secure (processes can directly access part of addr space of other processes)

**Message Passing:** processes want to exchange data send & receive msgs

One send and one receive:

```
send(addr, msg, length);
receive(addr, msg, length);
```

Adv: senders & receivers can be on diff machines; receiver can inspect msgs received before processing them

Disadv: hard to use (every data transfer requires send() & receive(), and receiving process must expect send())

**Defining Issues of Msg Passing:** Direct/Indirect communication, Blocking/Non-Blocking primitives, exception handling, quality of service

**Direct Comm:** send & rec calls specify processes as destination/source:

```
send(process, msg, length);
receive(process, msg, length);
```

No intermediary b/w sender & receiver. Ex. each phone hardwired to another

Proc exe rec call must know identity of all procs likely to send msgs (bad for servers)

**Indirect Comm:** send & rec primitives specify intermediary as destination or source (mailbox: sys obj created by kernel @ request of user process):

```
send(mailbox, msg, size);
receive(mailbox, msg, &size);
```

Diff proc can send msgs to same mailbox. A proc can rec msgs from procs it knows nothing about, can wait for msgs from diff senders (answ 1st msg rec)

**Private Mbox:** aka ports: proc requesting creation & children are only ones that can rec msgs through that mbox. Ceases to exist when proc requesting its creation (and all children) terminates.

**Public Mboxes:** owned by sys, & shared by all procs having right to rec msgs through it (works best when all procs on same machine). Survives termination of proc that requested creation. Ex. msg queues

**Blocking Primitives:** blocking send doesn't return til receiving process has received msg; no buffering needed. *Blocking receive* doesn't return til msgs have been received (std choice, but not good choice for direct comm)

**Non-Blocking Primitives:** *Non-blocking send* returns as soon as msg has been accepted for delivery by OS (assuming OS can store msg in a buffer, std choice). *Non-blocking receive* returns as soon as it has either retrieved msg or learned mbox is empty (retrieve()) acts as receive()

Simulating blocking receives: can sim blocking receive w/ non-blocking receiving within a loop (aka busy wait)

```
do {
    Code = receive(mbox, msg, size);
    sleep(1); // delay
```

```
} while (code == EMPTY_MBOX);
```

Simulating blocking sends: can sim blocking send w/ 2 non-blocking sends & a blocking receive

Sender sends msg & requests ACK -> sender wait for ACK from receiver using blocking receive -> receivers sends ACK

Non-blocking primitives require buffering to let OS store msgs that have been sent but not received. Buffers have a bounded capacity, but theoretically unlimited capacity.

**Exception Condition Handling:** must specify what to do if 1/2 processes dies. Esp important when the 2 procs on diff machines (must handle host failures & network partitions)

**Quality of Service:** when sender & receiver on diff machines, msgs can be lost, corrupted or duped, and can arrive out of sequence. Can still decide to provide reliable msg delivery

**Datagrams:** msgs sent individually (can be lost, duped, out of seq)

Reliable: msgs resent until ACK. Unreliable: msgs not ACK (works well when msg requests reply which acts as implicit ACK)

UDP (User Datagram Protocol) best known datagram protocol. Provides unreliable datagram services which is best for short interactions

**Virtual Circuits:** establishes logical connection b/w sender & receiver. Msgs guaranteed to arrive in seq w/o lost/duped msgs.

Requires virtual connection before sending any data. Best for transmitting large data amounts requiring sending several msgs (FTP, HTTP)

**Streams:** like virtual circuits, but does NOT preserve msg boundaries (seamless stream of bytes).

TCP (Transmission Control Protocol): best known stream protocol, providing reliable stream serv. Heavyweight (needs 3 msgs to estab virtual connection)

**Remote Procedure Calls (RPC):** applies to client-server model

```
send_req(args);
rcv_reqs(&args);
process(args, &results);
send_reply(results);
```

```
rcv_reply(&results);
```

Adv: hides all details of msg passing, provides higher level of abstraction, extends well-known model of programming

Disadv: illusion not perfect (RPCs don't behave exactly like regular procedure calls, client & server don't share same addr space), programmer must be aware of differences

User program contains user code & calls user stub that appears to call server procedure rpx(xyz, args, &results);

**User stub procedure generated by RPC package:** parcs args into request msg & performs required conversions (arg marshaling) -> sends request msg -> waits for server reply -> unpacks results & performs required conversions (arg unmarshaling)

Server stub: generic server generated by RPC pack waits for client requests, unpack request args & performs required data conversions, calls appropriate server procedure (which is written by user & does actual processing), packs results into reply msg (performs required conversions), sends reply msg

Client & server processes don't share same addr space: no global variables, can't pass by ref, can't pass dynam data structs via pointers; RPC can pass args by value & result (passes curr val to RP & copies returned val in user program).

## CHAPTER 5: CONCURRENCY

Multi App: invented to allow processing time to be shared among active apps

Structures Apps: extension of modular design & struct programming

OS Structure: OS themselves implemented as set of processes/threads

## Key Terms:

**Atomic Operation:** function/action implemented as sequence of >=1 instructions appearing indivisible. Sequence is guaranteed to exe as a group or not at all

**Critical Section:** section within proc that requires access to shared resources & must not be exe while another process is in a corresponding section of code

**Deadlock:** situation where >=2 processes unable to proceed bc each is waiting for one of the others to do something

**LiveLock:** >=2 procs continuously change states in response to changes in other proc(s) w/o doing useful work

**Mutual Exclusion:** requirement when one process is in crit sect that accesses shared resources, no other proc may be in crit sect that accesses any of those shared resources

Principles: Interleaving & overlapping (ex of concur processing), Uniproccessor - relative speed of exe of procs can't be predicted; dependent on activities of other procs, OS interrupt handling, OS scheduling policies

Difficulties: sharing of global resources, hard for OS to optimize resource allocation management, hard to locate programming errors

**Rare condition:** occurs when multi procs/threads read & write data items. Final result depends on order of exe ("loser" updates last, determines final val)

OS Concerns (Design & mgmt. issues): OS keeping track of various procs, allocation & deallocate resources for each active proc, protect data & phys resources of each proc against interference, ensure procs & outputs are independent of processing speed

## Interaction Process

Degree of Awareness	Relationship	Influence of 1 Process on the other	Potential Ctl Probs
Processes unaware of each other	Competition	- Results of 1 process independent of others - Timing may be affected	- Mut Exclusion (renewable) - Starvation
Processes indirectly aware of each other	Coop by sharing	- Results of 1 proc may depend on info fr others - Timing may be affected	- Mut Exclusion (renewable) - Starvation - Data coherence
Processes directly aware of each other (have comm prim available to them)	Coop by comm	- Results of 1 proc may depend on info from others - Timing may be affected	- Deadlock (consumable) - starvation

**Resource Competition:** concurr procs competing for use of same resource (I/O devices, mem, proc time, clock). Control Probs: need mutual exclusion, deadlock, starvation

**Mutual Exclusion Requirements:** Must be enforced, proc that halts has to w/o interfering w/ other procs, no deadlock/starvation, proc must not be denied access to crit sect when no other proc using it, no assume made about relative proc speeds/# of procs, proc remains inside crit sect for finite time only

**Mutual Exclusion hardware supp:** Special machine instructions - compare & swap (compare & exchange instruction). Compare made b/w mem val & test val. If same, swap.

## Compare & Swap Instructions:

```
const int n = /* # processes */;
int bolt;
void P(int i) {
    while (true) {
        while (compare_and_swap(bolt, 0, 1) == 1)
            /* do nothing */;
        /* crit section */;
        bolt = 0;
        /* remainder */;
    }
}

void main() {
    bolt = 0;
    parbegin (P(1), P(2), ..., P(n));
}
```

## Exchange Instructions:

```
const int n = /* # processes */;
int bolt;
void P(int i) {
    int keyi = 1;
    while (true) {
        do exchange (keyi, bolt)
        while (keyi != 0);
        /* crit section */;
        bolt = 0;
        /* remainder */;
    }
}

void main() {
    bolt = 0;
    parbegin (P(1), P(2), ..., P(n));
}
```

**Special Machine Instruct Adv:** applicable to any # of pros on single/multiple processors sharing main mem, simple & easy to verify, can be used to supp multi crit sects (each can be defined by its own variable)

**Disadv:** busy-wait employed, while a proc is waiting for access to crit section it cont to consume processor time, starvation is possible when proc leaves crit sect & >= 1 proc is waiting, deadlock is possible

Common Concurrency Mechanisms	
<b>Semaphore</b>	An integer value used for signaling among processes. Only three operations may be performed on a semaphore, all of which are atomic: initialize, decrement, and increment. The decrement operation may result in the blocking of a process, and the increment operation may result in the unblocking of a process. Also known as a <b>counting semaphore</b> or a <b>general semaphore</b>
<b>Binary Semaphore</b>	A semaphore that takes on only the values 0 and 1.
<b>Mutex</b>	Similar to a binary semaphore. A key difference between the two is that the process that locks the mutex (sets the value to zero) must be the one to unlock it (sets the value to 1).
<b>Condition Variable</b>	A data type that is used to block a process or thread until a particular condition is true.
<b>Monitor</b>	A programming language construct that encapsulates variables, access procedures and initialization code within an abstract data type. The monitor's variable may only be accessed via its access procedures and only one process may be actively accessing the monitor at any one time. The access procedures are <i>critical sections</i> . A monitor may have a queue of processes that are waiting to access it.
<b>Event Flags</b>	A memory word used as a synchronization mechanism. Application code may associate a different event with each bit in a flag. A thread can wait for either a single event or a combination of events by checking one or multiple bits in the corresponding flag. The thread is blocked until all of the required bits are set (AND) or until at least one of the bits is set (OR).
<b>Mailboxes/Messages</b>	A means for two processes to exchange information and that may be used for synchronization.
<b>Spinlocks</b>	Mutual exclusion mechanism in which a process executes in an infinite loop waiting for the value of a lock variable to indicate availability.

**Semaphore:** A variable w/ int val where only 3 ops are defined -> no way to inspect/manip semaphores other than those ops: 1) may be init to nonneg int val 2) semWait() on decrements val 3) semSignal() on increments val

Cons: no way to know before proc decrements semaphore if block, no way to know which proc will cont immediately on uniproc sys when 2 running concurr, don't know if another proc is waiting so # of unblocked proc is 0 or 1

## Semaphore Primitives

```
struct semaphore {
    int count;
    queueType queue;
};

void semWait(semaphore s) {
    s.count--;
    if (s.count < 0) {
        /* place process in s.queue */;
        /* block this process */;
    }

    void semSignal(semaphore s) {
        s.count++;
        if (s.count <= 0) {
            /* remove process P from s.queue */;
            /* place process P on ready list */;
        }
    }
}
```

**Strong semaphores:** the proc that's been blocked the longest released fr queue (FIFO)

**Weak semaphores:** order where procs removed from queue not specified

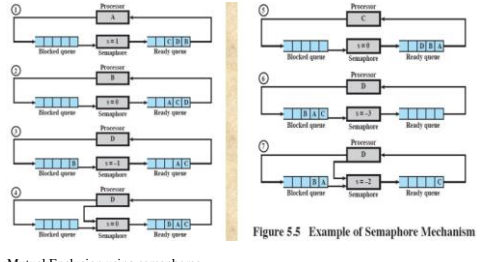


Figure 5.5 Example of Semaphore Mechanism

**Mutual Exclusion using semaphores**

```
const int n = /* # processes */;
semaphore s = 1;
void P(int i) {
    while (true) {
        semWait(s);
        /* crit section */;
        semSignal(s);
        /* remainder */;
    }
}

void main() {
    parbegin (P(1), P(2), ..., P(n));
}
```

**Producer/Consumer Problem:** Ensure that produced can't add data into full buff & consumer can't remove data from empty buff

General situation: >= 1 producers gen data & placing in buff where a consumer taking items out individually, but only 1 producer/consumer may access buff at any one time.

**Solution using semaphores:**

```
/* program producerconsumer */
semaphore n = 0, s = 1;
void producer() {
    while (true) {
        produce();
        semWait(s);
        append();
        semSignal(s);
        semSignal(n);
    }
}

void consumer() {
    while (true) {
        semWait(n);
        semWait(s);
        take();
        semSignal(s);
        consume();
    }
}

void main() {
    parbegin (producer, consumer);
}
```

**Implementation of semaphore:** imperative semWait & semSignal ops implemented as atomic primitives in hardware/firmware, Dekker's/peterson's algorithm can be used. Use one of hardware-supported schemes for mutual exclusion

**Monitors:** construct that provides equiv functionality to semaphores that is easier to control that is implemented in a number of langs, & as program lib. Software module consisting of >= 1 procedures, an init sequence, & local data

**Characteristics:** local data variables accessible only by mon's procedures (no ext procedure), process enters mon by invoking one of its procedures, only 1 proc exe in mon at a time.

**Synchronization:** achieved by condition variables contained within & only accessible within mon. cwait(c) suspends exe of calling proc on condition c, csignal(c) resumes exe of some proc blocked after cwait on same condition

**Solution to bounded-buff prod/consum using monitor**

<pre>/* program producerconsumer */ monitor boundedbuff {     char buf[N];     int nextIn, nextOut;     bool full, empty;     conditionVariable fullCondition, emptyCondition;      void append (char c) {         if (count == N) wait(emptyCondition);         if (full) wait(fullCondition);         buf[nextIn] = c;         nextIn = (nextIn + 1) % N;         count++;         if (count == N) signal(emptyCondition);     }      void take (char c) {         if (count == 0) wait(fullCondition);         if (nextOut == nextIn) wait(emptyCondition);         c = buf[nextOut];         nextOut = (nextOut + 1) % N;         count--;         if (count == 0) signal(fullCondition);     } }</pre>	<pre>void producer() {     char x;     while (true) {         produce(x);         append(x);     } }  void consumer() {     char x;     while (true) {         take(x);         consume(x);     } }  void main() {     parbegin (producer, consumer); }</pre>
---	---

**Message Passing:** when proc interact with another these requirements must be satisfied and msg: synchronization (ensure mut exclu) and comm (to exchange info)

**Blocking send, Blocking rec:** both sender & rec blocked til msg delivered, allowing for tight synchronization b/w procs

**Nonblocking send, blocking rec:** sender continues but rec blocked til requested msg arrives. Most useful combo. Sends >= 1 msg to a variety destinations asap

**Nonblocking send, nonblocking rec:** neither party required to wait

**Readers/Writers Prob:** data area shared among many procs, following conditions must be satisfied: any # of readers may read file simultaneously, only 1 writer may write at a time, if writer is writing, no reader may read it

